




Figure 1: **LAYOUT TEMPLATE** — not yet a result. Intended structure of the indoor figure: one instance, consecutive observed frames, the per-frame open-vocabulary prediction above and our emitted label below. The number of label runs shown matches the measured switch count for the candidate instance; the class names are deliberately left as ? because no stored run persists the per-frame label sequence, and recovering it requires a replay that has not been run. Candidate instances, with measured switch count, track length and mask IoU, are listed in the accompanying shortlist; the chosen one fills this layout unchanged.